

Promineo Tech Lesson - Week 8

 promineotech.openclass.ai/resource/lesson-659c2dc3fe775297cad773cb

MySQL Part 3

This lesson contains Questions 21-30 (Part 3) of the OpenClass MySQL Lesson, and contains five (5) coding questions, each followed by a mastery question.

Vocabulary Reminder

- **Database:** A structured collection of data stored and organized for efficient retrieval and manipulation.
- **Schema:** A blueprint or structure that defines the logical organization and relationships of database objects, such as tables, views, and constraints.
- **DBMS** (Database Management System): Software that manages the storage, retrieval, and manipulation of data in a database.
- **RDBMS** (Relational Database Management System): An RDBMS organizes data based on the relational model, consisting of tables with rows and columns.
- **SQL** (Structured Query Language): is a language that is used by an RDBMS to interact with and manage relational data. SQL is a Standardized Language.
- **Query:** A request for data retrieval or manipulation from a database using a structured query language (SQL).
- **Primary Key:** A unique identifier for each record (row) in a table, used to ensure data integrity and facilitate record retrieval.
- **Foreign Key:** A field in one table that refers to the primary key of another table, establishing a relationship between the two tables.
- **Table:** A collection of related data organized in rows and columns in a relational database.
- **Entity:** A distinct object or concept in the real world that is represented in a database table.
- **Index:** A data structure that improves the speed of data retrieval operations by enabling efficient searching and sorting.
- **Attribute:** A characteristic or property of an entity that is stored as a column in a database table.
- **Transaction:** A logical unit of work that consists of one or more database operations, which must be performed atomically and consistently.
- **Normalization:** The process of organizing data in a database to eliminate redundancy and dependency issues.
- **ACID** (Atomicity, Consistency, Isolation, Durability): A set of properties that ensure reliability and consistency in database transactions.
- **Data Integrity:** The accuracy, consistency, and reliability of data stored in a database.
- **Query Optimization:** The process of selecting the most efficient execution plan for a database query to improve performance.
- **Data Warehousing:** The process of collecting, organizing, and storing large volumes of data from various sources for analysis and reporting.
- **Data Mining:** The process of discovering patterns, relationships, and insights from large datasets using statistical and machine learning techniques.
- **Backup and Recovery:** The process of creating backups of database data and implementing strategies to restore data in case of system failures or data loss.
- **CRUD** (Create, Read, Update, and Delete): The Operations that can be performed on a DBMS or RDBMS.
- **DML** (Data Manipulation Language): The language keywords that help manage and manipulate data in the database.
- **Examples of DML:** *SELECT, INSERT, UPDATE, and DELETE*
- **DDL** (Data Definition Language): The language keywords that help to define the structure or schema of the database
- **Examples of DDL:** *CREATE, ALTER, and DROP*

CRUD Operations in SQL

When we look at the **CRUD** operations on the data in the database, **CRUD** operations are all **DML** statements, as follows:

- **Create:** SQL *INSERT* statement
- **Read:** SQL *SELECT* statement
- **Update:** SQL *UPDATE* statement
- **Delete:** SQL *DELETE* statement

SQL Syntax

INSERT

```

INSERT [INTO] tbl_name
    [ (col_name [, col_name] ...)]
    { VALUES (value_list) [, value_list] ] ... };
value:
    { expr | DEFAULT}
value_list:
    value [, value] ...
assignment:
    col_name =
        value
        | [row_alias.]col_name
        | [tbl_name.]col_name
        | [row_alias.]col_alias
assignment_list:
    assignment [, assignment] ...

```

SELECT

```

SELECT select_expr [, select_expr] . . .
    [FROM table_references]
    [JOIN table_references { ON (col_name) | USING (col_name = col_name) } ]
    [WHERE where_condition]
    [GROUP BY {col_name | expr | position}, ... ]
    [ORDER BY {col_name | expr | position}
        [ASC | DESC], ... ]
    [LIMIT row_count];

```

UPDATE

```

UPDATE table_reference
SET assignment_list
    [WHERE where_condition]
    [ORDER BY ...]
    [LIMIT row_count];
value:
    { expr | DEFAULT}
assignment:
    col_name = value
assignment_list:
    assignment [, assignment] ...

```

DELETE

```

DELETE FROM tbl_name
    [WHERE where_condition]
    [ORDER BY ...]
    [LIMIT row_count];

```

References

Sakila Database

For all of our SQL lessons, we are going to use the **Sakila** Database. Each id column which is named *tableName_id* is a **PRIMARY KEY** in that table. Notice that some of those columns are used in subsequent tables as well, in those tables, that first *tableName_id* would be stored as a **FOREIGN KEY** in the subsequent table.

The **Table** and **Column Names** in this database are these:

Table Name	Column Names
actor	actor_id, first_name, last_name last_update
address	address_id, address, address2, district, city_id (FK), postal_code, phone, location, last_update
category	category_id, name, last_update
city	city_id, city, country_id (FK), last_update
country	country_id, country, last_update
customer	customer_id, store_id (FK), first_name, last_name, email, address_id (FK), active, create_date, last_update
film	film_id, title, description, release_year, language_id (FK), original_language_id (FK), rental_duration, rental_rate, length, replacement_cost, rating, special_features, last_update
film_actor	actor_id (FK), film_id (FK), last_update
film_category	film_id (FK), category_id (FK), last_update
film_text	film_id (FK), title, description
inventory	inventory_id, film_id, store_id (FK), last_update
language	language_id, name, last_update
payment	payment_id, customer_id (FK), staff_id (FK), rental_id (FK), amount, payment_date, last_update
rental	rental_id, rental_date, inventory_id (FK), customer_id (FK), return_date, staff_id (FK), last_update
staff	staff_id, first_name, last_name, address_id (FK), picture, email, store_id (FK), active, username, password, last_update
store	store_id, manager_staff_id (FK), address_id (FK), last_update

Each OpenClass question has a database attached to it. The only requirements necessary in the Solution Box are the SQL statements that accomplish what is being requested.

Reference: Sakila Database

21. *SELECT* Statement -- *JOIN* two tables, *payment* and *customer* -- Use column aliases

Retrieve the *customer_id*, the customer's first and last names, and the average of the *amount* paid for a rental in the *payment* table per customer, rounded to two (2) decimal places. Limit your results to the first 5 rows. Use the following column aliases in your query:

- *customer_id* --> "Id"
- *first_name* --> "First Name"
- *last_name* --> "Last Name"
- rounded average --> "Average Spent"

Remember, when using an aggregate function, a *GROUP BY* clause may be important.

Sample Test Case #1

Expected STDOUT

Id	First Name	Last Name	Average Spent
----	------------	-----------	---------------

1	MARY	SMITH	3.71
2	PATRICIA	JOHNSON	4.77
3	LINDA	WILLIAMS	5.21
4	BARBARA	JONES	3.72
5	ELIZABETH	BROWN	3.81

CREATE TABLE Syntax

When creating a database, this statement will allow a change to be made to the structure of the database or schema. This is a **DDL** statement.

```
CREATE [TEMPORARY] TABLE [IF NOT EXISTS] tbl_name
    (create_definition,...)
    [table_options]
    [partition_options]

create_definition: {
    col_name column_definition
  | {INDEX | KEY} [index_name] [index_type] (key_part,...)
    [index_option] ...
  | {FULLTEXT | SPATIAL} [INDEX | KEY] [index_name] (key_part,...)
    [index_option] ...
  | [CONSTRAINT [symbol]] PRIMARY KEY
    [index_type] (key_part,...)
    [index_option] ...
  | [CONSTRAINT [symbol]] UNIQUE [INDEX | KEY]
    [index_name] [index_type] (key_part,...)
    [index_option] ...
  | [CONSTRAINT [symbol]] FOREIGN KEY
    [index_name] (col_name,...)
    reference_definition
  | check_constraint_definition
}
```

```
column_definition: {
    data_type [NOT NULL | NULL] [DEFAULT {literal | (expr)} ]
    [AUTO_INCREMENT] [UNIQUE [KEY]] [[PRIMARY] KEY]
    [reference_definition]
  | data_type
    [VIRTUAL | STORED] [NOT NULL | NULL]
    [UNIQUE [KEY]] [[PRIMARY] KEY]
    [reference_definition]
}
```

CREATE TABLE **employee** Example

```
CREATE TABLE employee (
    employee_id SMALLINT unsigned NOT NULL,
    first_name VARCHAR(45) NOT NULL,
    last_name VARCHAR(45) NOT NULL,
    address_id SMALLINT unsigned NOT NULL,
    email VARCHAR(50) DEFAULT NULL,
    store_id TINYINT unsigned NOT NULL,
    employee_level TEXT CHECK( employee_level IN ('FAIR','GOOD','GREAT','EXCELLENT') ) ,
    title VARCHAR(25) DEFAULT NULL,
    PRIMARY KEY (employee_id),
    FOREIGN KEY (address_id) REFERENCES address (address_id),
    FOREIGN KEY (store_id) REFERENCES store (store_id)
);
```

References:

- [MySQL CREATE TABLE Documentation](#)
- [MySQL Data Types Documentation](#)

22. CREATE TABLE rewards

Create a new table: rewards

Add the following fields:

- `rewards_id` SMALLINT
- `customer_id` SMALLINT
- `status` VARCHAR(20) DEFAULT 'MEMBER'
- `discount_percent` DOUBLE DEFAULT 0.0
- `year_joined` YEAR DEFAULT CURRENT_YEAR

All fields need to be `NOT NULL` or use the `DEFAULT` keyword

Add the following PRIMARY and FOREIGN KEYS

- PRIMARY KEY (`rewards_id`)
- FOREIGN KEY (`customer_id`) REFERENCES customer(`customer_id`)
- FOREIGN KEY (`status`) REFERENCES reward_status (`status`)

Test the creation by doing a retrieval of all data from this table -- which will show the column names.

`SELECT * FROM rewards;` will retrieve all of the data from this table (or in this case, show the column names).

Sample Test Case #1

Expected STDOUT

```
rewards_id  customer_id  status  discount_percent  year_joined
```

status Table

This table has been created in your database, and the table is called `status`. This is one way to store information in the database, and restrict values for a particular field.

There is a new table added into this database called `status`, which has the following columns and values:

Table Name	Column Names	Allowed Values
status	status	'MEMBER', 'SILVER', 'GOLD', 'PLATINUM'

23. INSERT INTO rewards table

Add the following 4 rows into the `rewards` table:

- (1, 1, 'PLATINUM', 0.20, '2000')
- (2, 2, 'GOLD', 0.15, '2010')
- (3, 3, 'SILVER', 0.10, '2015')
- (4, 4, 'MEMBER', 0.05, '2020')

Test those inserts by retrieving all information from the `rewards` table.

Sample Test Case #1

Expected STDOUT

rewards_id	customer_id	status	discount_percent	year_joined
1	1	PLATINUM	0.2	2000
2	2	GOLD	0.15	2010
3	3	SILVER	0.1	2015
4	4	MEMBER	0.05	2020

Sakila customer table with columns

Table Name	Column Names
customer	customer_id, store_id (FK), first_name, last_name, email, address_id (FK), active, create_date, last_update
rewards	rewards_id, customer_id, status, discount_percent, year_joined

24. **INSERT** a new customer in the **customer** table

Insert customer "Mary Mallows" with the following data:

- **customer_id**: 1000
- **store_id**: 1
- **first_name**: "Mary"
- **last_name**: "Mallows"
- **email**: "mm@gmail.com"
- **address_id**: 1
- **active**: 1
- **create_date**: "2023-05-31 14:26:19"
- **last_update**: "2023-05-31 14:26:19"

Add the following row into the **rewards** table: (1000, 1000, 'PLATINUM', 0.20, '2023')

Test those inserts by retrieving this **customer_id** from **customer** and **rewards**

Use a **JOIN** in your **SELECT** statement.

SELECT * FROM customer

JOIN rewards USING (customer_id)

WHERE customer_id = 1000;

Sample Test Case #1

Expected STDOUT

customer_id	store_id	first_name	last_name	email	address_id	active	create_date	last_update	rewards_id	status
1000	1	Mary	Mallows	mm@gmail.com	1	1	2023-05-31 14:26:19	2023-05-31 14:26:19	1000	PLAT

25. **SELECT** customer "Mary Mallows" with **rewards** data

Retrieve the first name & last name from the customer with the name "Mary Mallows" along with her **rewards** information.

SELECT first_name, last_name, rewards.* FROM customer

Then use a **JOIN** in your **SELECT** statement to get the rest of the information.

Sample Test Case #1

Expected STDOUT

first_name	last_name	rewards_id	customer_id	status	discount_percent	year_joined
Mary	Mallows	1000	1000	PLATINUM	0.2	2023

26. **UPDATE** customer "Mary Mallows"

"Mary Mallows" got married this week. You need to change the following:

- Her last name to "Smith".
- Her e-mail address to "ms@gmail.com".
- Don't forget to update the **last_update** column to "2023-06-19 23:00:00".

The **customer_id** of "Mary Mallows" is 1000.

SELECT the customer record showing the updated information.

Note: In real life, we would set `last_update` to `CURRENT_TIMESTAMP`, but for testing purposes, we will set it to the same value as the comparison!

Sample Test Case #1

Expected STDOUT

customer_id	store_id	first_name	last_name	email	address_id	active	create_date	last_update
1000	1	Mary	Smith	ms@gmail.com	1	1	2023-05-31 14:26:19	2023-06-19 23:00:00

27. *DELETE customer "Mary Smith"*

Do not forget the WHERE clause in an SQL DELETE statement!

SELECT the customer record from **customer** showing that it has been deleted.

Sample Test Case #1

Expected STDOUT

customer_id	store_id	first_name	last_name	email	address_id	active	create_date	last_update
-------------	----------	------------	-----------	-------	------------	--------	-------------	-------------

ORDER BY Clause

When retrieving data from a relational database system, it is sometimes important to order that data in a particular way. The **ORDER BY** clause allows the data to be displayed in a particular order. The **ORDER BY** clause is an optional part in a **SELECT** statement, and is used to sort the result set in ascending (**ASC**) or descending (**DESC**) order by whatever **column_name**, expression or position is chosen. If omitted, the data will be displayed in the order that it is retrieved.

The **ORDER BY** keyword sorts the records in ascending order by default. To sort the records in descending order, use the **DESC** keyword. The **ASC** keyword is also used for clarity, even though records are sorted in ascending order by default.

ORDER BY Syntax

```
[ ORDER BY { column_name | expression | position } [ DESC | ASC ] ... ]
```

28. *SELECT from film*

List the bottom 5 films with the lowest rental count, retrieving the title, and the rental count AS "rental_count". Order the results from lowest to highest rental count, and Limit the result to 5.

This query requires information from three (3) tables: **film**, **inventory** and **rental**.

Additionally, use the **GROUP BY title**, because of the aggregate function to achieve the COUNT, and an **ORDER BY** to correctly order your results.

Sample Test Case #1

Expected STDOUT

title	rental_count
HARDLY ROBBERS	4
MIXED DOORS	4
TRAIN BUNCH	4
BRAVEHEART HUMAN	5
BUNCH MINDS	5

JOIN Clause

A **JOIN** clause is part of the **SELECT** statement in SQL, and is used to combine rows from two (2) or more tables, based on a related column between the tables. The concept used here connects a **FOREIGN KEY** (FK) in one table to a **PRIMARY KEY** (PK) in another table.

When joining tables in a relational database management system, there are four distinct types of joins:

- **(INNER) JOIN**: An **inner join** returns records that have matching values in both tables.
- **RIGHT (OUTER) JOIN**: A **right (outer) join** returns **all records** from the **right** table, and the **matched records** from the **left** table.
- **LEFT (OUTER) JOIN**: A **left (outer) join** returns **all records** from the **left** table, and the **matched records** from the **right** table.
- **FULL (OUTER) JOIN**: A **full (outer) join** returns **all records** from **both** tables when there is a match in either the left or right table.

In most cases, a **SELECT** is looking for the rows that have matching values in both tables, so a **INNER JOIN** is used.

USING VS. ON

When joining tables, there are two types of syntax that can be used. If the **PK/FK** pair of columns have **exactly the same name in each table**, then the keyword **USING** can be used. If the **PK/FK** pair has a **different name in each table**, then the **ON** clause can be used. Imagine joining these two tables from the **Sakila** database.

Table Name	Column Names
film	film_id, title, description, release_year, language_id (FK), original_language_id (FK), rental_duration, rental_rate, length, replacement_cost, rating, special_features, last_update
language	language_id (PK), name, last_update

There are two Examples below, both joining the **film** table with the **language** table on the **language.language_id** **PK** column.

Example 1: Join **film** and **language** using **film.language_id** to **language.language_id**

Notice that in this example, the column names are exactly the same in each table, so the **USING** keyword can be used.

```
SELECT * from film
INNER JOIN language USING (language_id);
```

The **ON** keyword can also be used, and the following example will do the exact same thing as the above version:

```
SELECT * from film
INNER JOIN language ON language.language_id = film.language_id;
```

Example 2: Join **film** and **language** using **film.original_language_id** to **language.language_id**

Notice that in this example, the column names are **different** in each table, so the **ON** keyword **must** be used.

```
SELECT * from film
INNER JOIN language ON film.original_language_id = language.language_id;
```

This query can also be written using **table aliases**:

```
SELECT * from film f
INNER JOIN language l ON f.original_language_id = l.language_id;
```

29. **SELECT** from **film**

- List the top 5 film titles with the highest rental count
- Retrieve the rental count AS "rental_count"
- Remember that COUNT() is an aggregate.
- Order the results from highest to lowest rental count
- Limit the result to 5.

This query requires information from three (3) tables: **film**, **inventory** and **rental**.

Don't forget the **GROUP BY title**, because of the aggregate function to achieve the COUNT, and an **ORDER BY** to correctly order your results.

Sample Test Case #1

Expected STDOUT

title	rental_count
BUCKET BROTHERHOOD	34
ROCKETEER MOTHER	33
SCALAWAG DUCK	32
RIDGEMONT SUBMARINE	32
JUGGLER HARDLY	32

30. *SELECT from film*

- Given the top 5 films with the highest rental count, ordered from highest to lowest by "rental_count". See *question 29 solution*.
- Retrieve the amount of money that these 5 top films have made AS "rental_amount"

NOTE: The **JOIN** in question 29. will give you all of the information that you need.

This query requires information from three (3) tables: **film**, **inventory** and **rental**.

Use **SUM()** to calculate the amount of money made, which is found by multiplying **rental_duration** by **rental_rate**.

Sample Test Case #1

Expected STDOUT

title	rental_count	rental_amount
BUCKET BROTHERHOOD	34	1187.6199999999997
ROCKETEER MOTHER	33	98.00999999999998
SCALAWAG DUCK	32	958.08000000000008
RIDGEMONT SUBMARINE	32	95.03999999999998
JUGGLER HARDLY	32	126.71999999999991